

Assignment 1

Please enter your **name**, **surname** and **student number** instead of "NAME-HERE" , "SURNAME-HERE" , "NUMBER-HERE" below.

In [112...

```
student = {  
    'name' : "Defne" ,  
    'surname' : "Odabaşı",  
    'studentNumber' : "2443604"  
}  
  
print(student)
```

```
{'name': 'Defne', 'surname': 'Odabaşı', 'studentNumber': '2443604'}
```

Part I: Classification Problem

1. The "IBM HR Analytics Employee Attrition Dataset" should be downloaded from the Kaggle website: <https://www.kaggle.com/datasets/pavansubhasht/ibm-hr-analytics-attrition-dataset>
2. Use a seed value of 12345 for random number generators to ensure reproducibility in your experiments. [Mandatory, 0 points]
3. Conduct Exploratory Data Analysis (EDA) to gain insights into the dataset characteristics. Employ statistical summaries and visualizations to uncover patterns and anomalies. [10 points]
4. Execute data preprocessing to enhance model performance if deemed necessary. This may include handling missing values, encoding categorical variables, feature scaling, and any other technique that could improve the results. [5 points]
5. Implement 5-Fold Cross Validation to assess the robustness of your models. This approach ensures that the evaluation of your model is as accurate as possible. [5 points]
6. Develop and evaluate models using K-Nearest Neighbors (KNN), Naive Bayes, Perceptron, and Logistic Regression algorithms. Document the performance of each model. [30 points]
7. Investigate the outcomes using appropriate metrics such as accuracy, precision, recall, F1 score, and ROC-AUC curve where applicable. [5 points]
8. Discuss the results. Reflect on which model yielded the best performance and hypothesize why this might be the case. Consider the algorithm's suitability for the data distribution, complexity, and balance of the dataset. [15 points]

Your Discussion Here (You can double click and edit this block)

In [113...

```
import pandas as pd  
import numpy as np  
import seaborn as sns  
import matplotlib.pyplot as plt  
from sklearn.model_selection import train_test_split, KFold  
from sklearn.preprocessing import StandardScaler  
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.linear_model import Perceptron
```

In [114...

```
seed_value = 12345 #seed value
np.random.seed(seed_value)
```

Exploratory Data Analysis (EDA)

Gaining insides into dataset characteristics. Employing statistical summaries and visualizations to uncover patterns and anomalies:

In [115...

```
# Your code here (you can add more blocks as you need).


employee_attrition = pd.read_csv("C:\\Users\\defne\\Desktop\\2023-2024FallSemester\\
employee_attrition.head()

# Please add comments where you think necessary.
```

Out[115]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Educa
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life
4	27	No	Travel_Rarely	591	Research & Development	2	1	

5 rows × 35 columns



In [116...

```
employee_attrition.shape
```

Out[116]:

(1470, 35)

In [117...

```
employee_attrition.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                    1470 non-null   int64
1   Attrition                            1470 non-null   object
2   BusinessTravel                       1470 non-null   object
3   DailyRate                            1470 non-null   int64
4   Department                           1470 non-null   object
5   DistanceFromHome                     1470 non-null   int64
6   Education                            1470 non-null   int64
7   EducationField                       1470 non-null   object
8   EmployeeCount                        1470 non-null   int64
9   EmployeeNumber                       1470 non-null   int64
10  EnvironmentSatisfaction               1470 non-null   int64
11  Gender                               1470 non-null   object
12  HourlyRate                           1470 non-null   int64
13  JobInvolvement                       1470 non-null   int64
14  JobLevel                             1470 non-null   int64
15  JobRole                              1470 non-null   object
16  JobSatisfaction                      1470 non-null   int64
17  MaritalStatus                       1470 non-null   object
18  MonthlyIncome                       1470 non-null   int64
19  MonthlyRate                          1470 non-null   int64
20  NumCompaniesWorked                  1470 non-null   int64
21  Over18                              1470 non-null   object
22  OverTime                             1470 non-null   object
23  PercentSalaryHike                   1470 non-null   int64
24  PerformanceRating                   1470 non-null   int64
25  RelationshipSatisfaction             1470 non-null   int64
26  StandardHours                       1470 non-null   int64
27  StockOptionLevel                    1470 non-null   int64
28  TotalWorkingYears                   1470 non-null   int64
29  TrainingTimesLastYear               1470 non-null   int64
30  WorkLifeBalance                     1470 non-null   int64
31  YearsAtCompany                      1470 non-null   int64
32  YearsInCurrentRole                  1470 non-null   int64
33  YearsSinceLastPromotion              1470 non-null   int64
34  YearsWithCurrManager                 1470 non-null   int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB

```

In [118...

```
employee_attrition.describe() #statistics from the data
```

Out[118]:

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNuml
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	1470.0000
mean	36.923810	802.485714	9.192517	2.912925	1.0	1024.8653
std	9.135373	403.509100	8.106864	1.024165	0.0	602.0243
min	18.000000	102.000000	1.000000	1.000000	1.0	1.000000
25%	30.000000	465.000000	2.000000	2.000000	1.0	491.2500
50%	36.000000	802.000000	7.000000	3.000000	1.0	1020.5000
75%	43.000000	1157.000000	14.000000	4.000000	1.0	1555.7500
max	60.000000	1499.000000	29.000000	5.000000	1.0	2068.0000

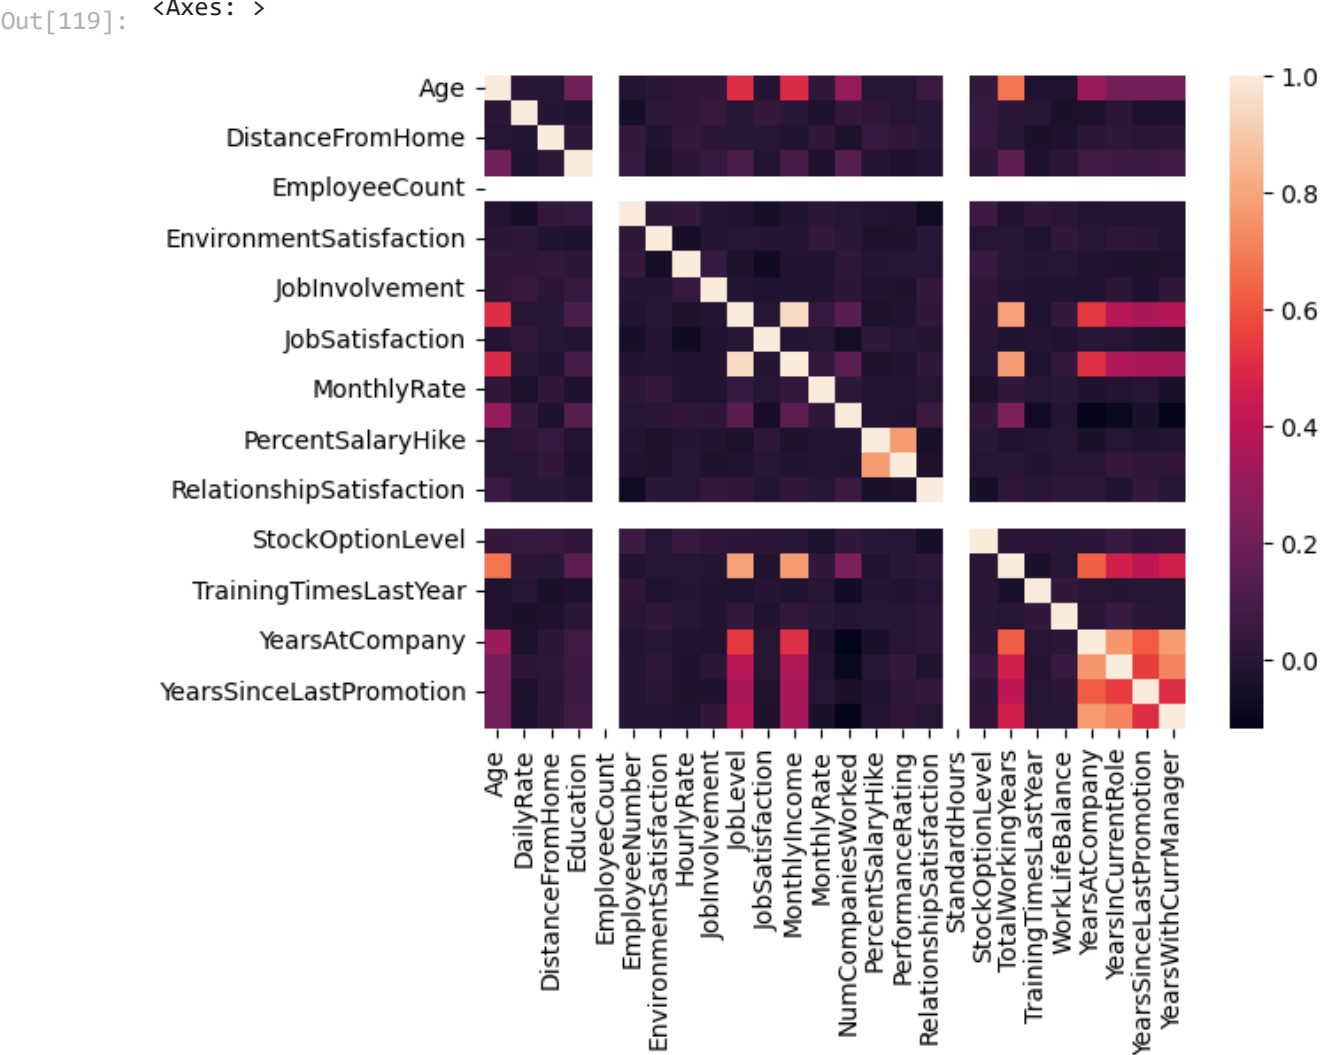
8 rows × 26 columns

In [119...

```
#Correlation plot
#after encoding it is provided again
sns.heatmap(employee_attrition.corr())
```

C:\Users\defne\AppData\Local\Temp\ipykernel_10008\4079675858.py:3: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

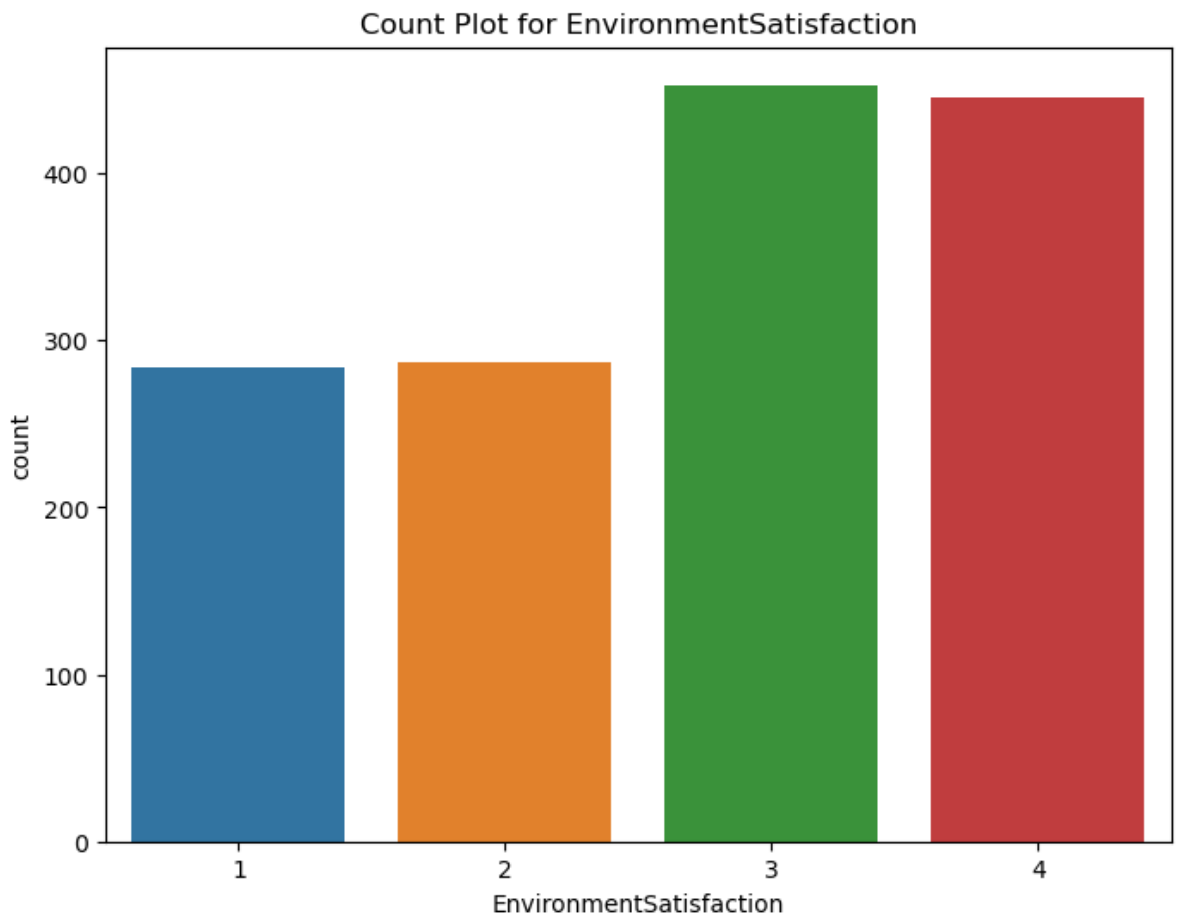
```
sns.heatmap(employee_attrition.corr())
```

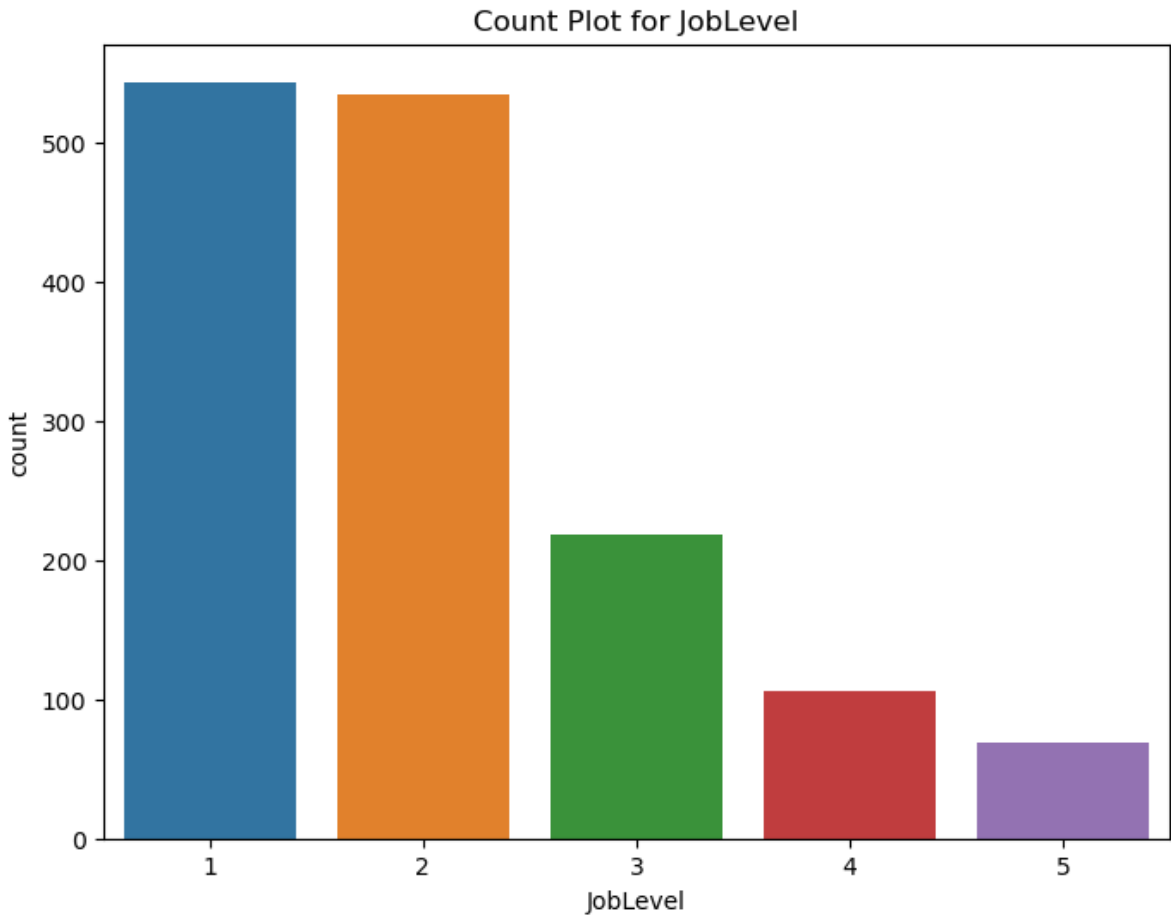
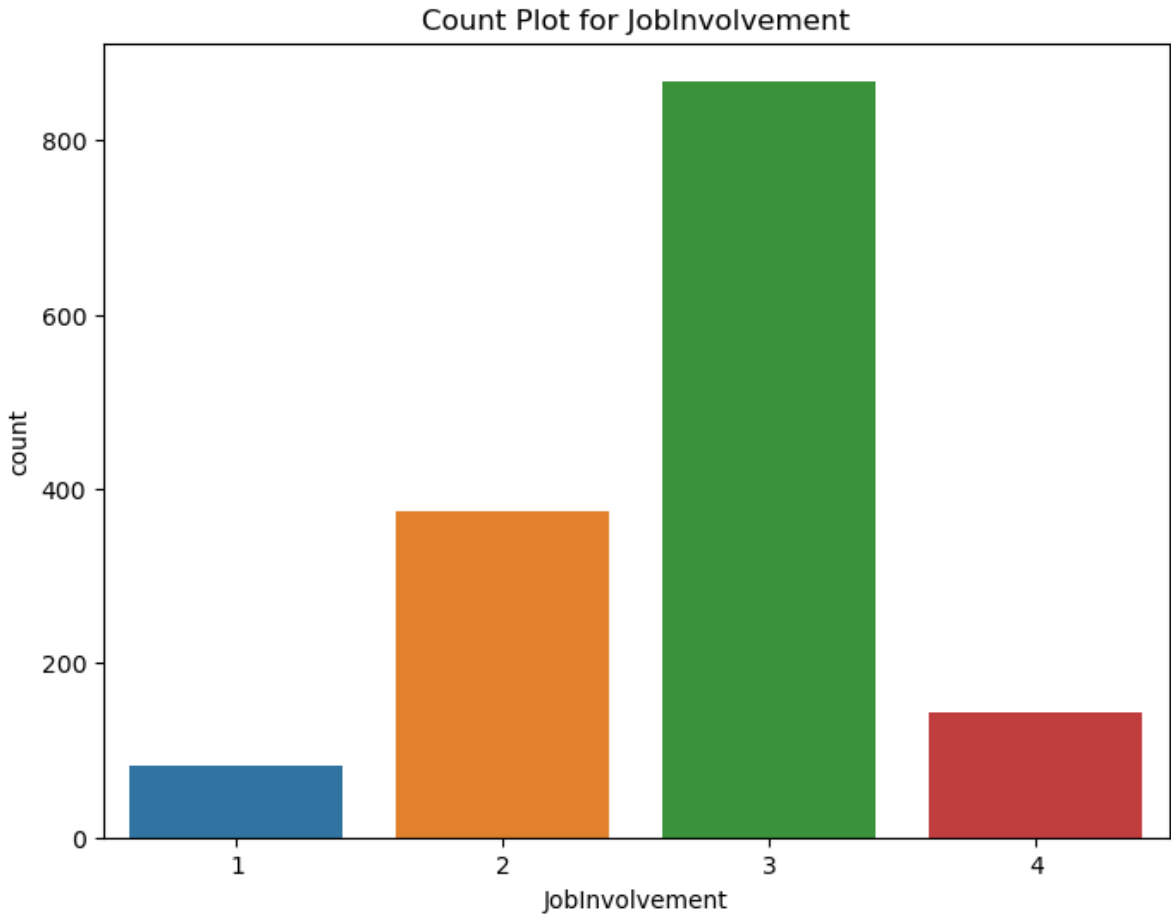


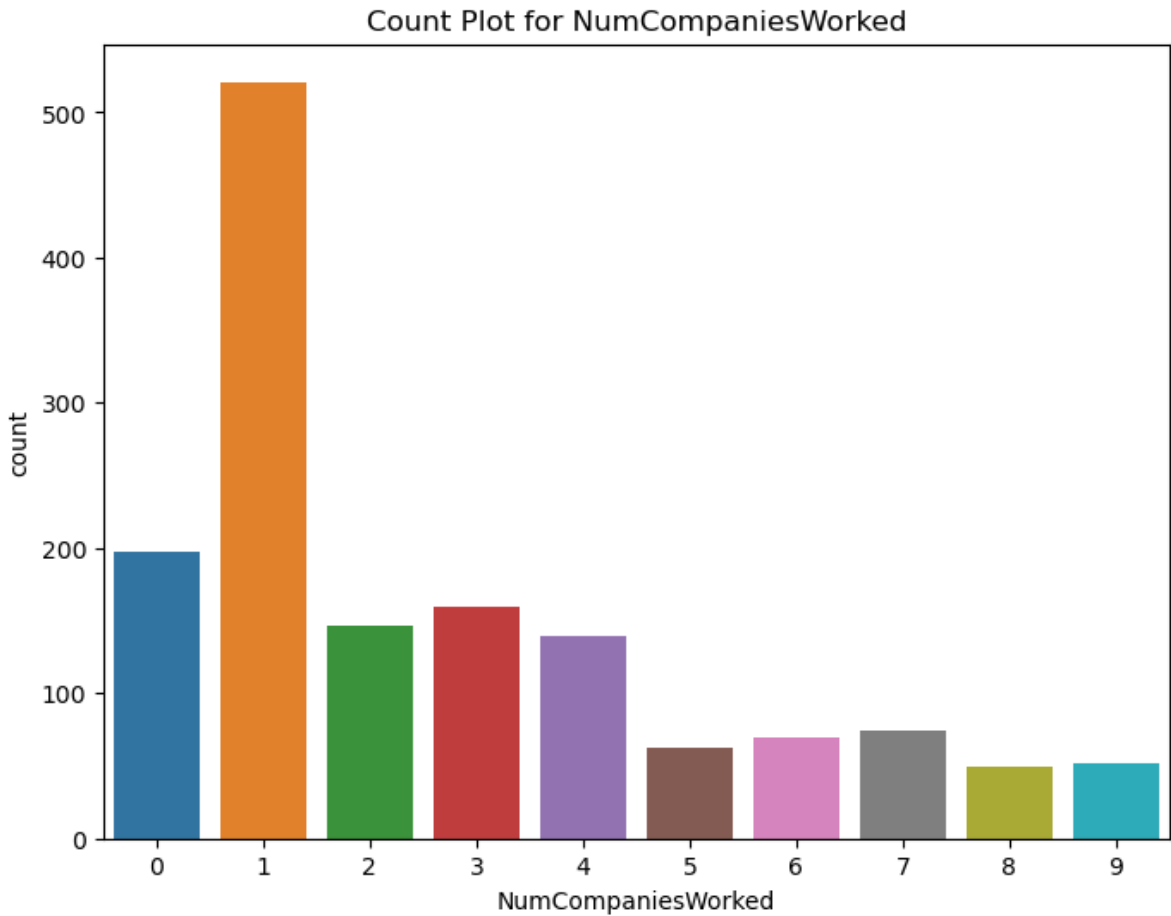
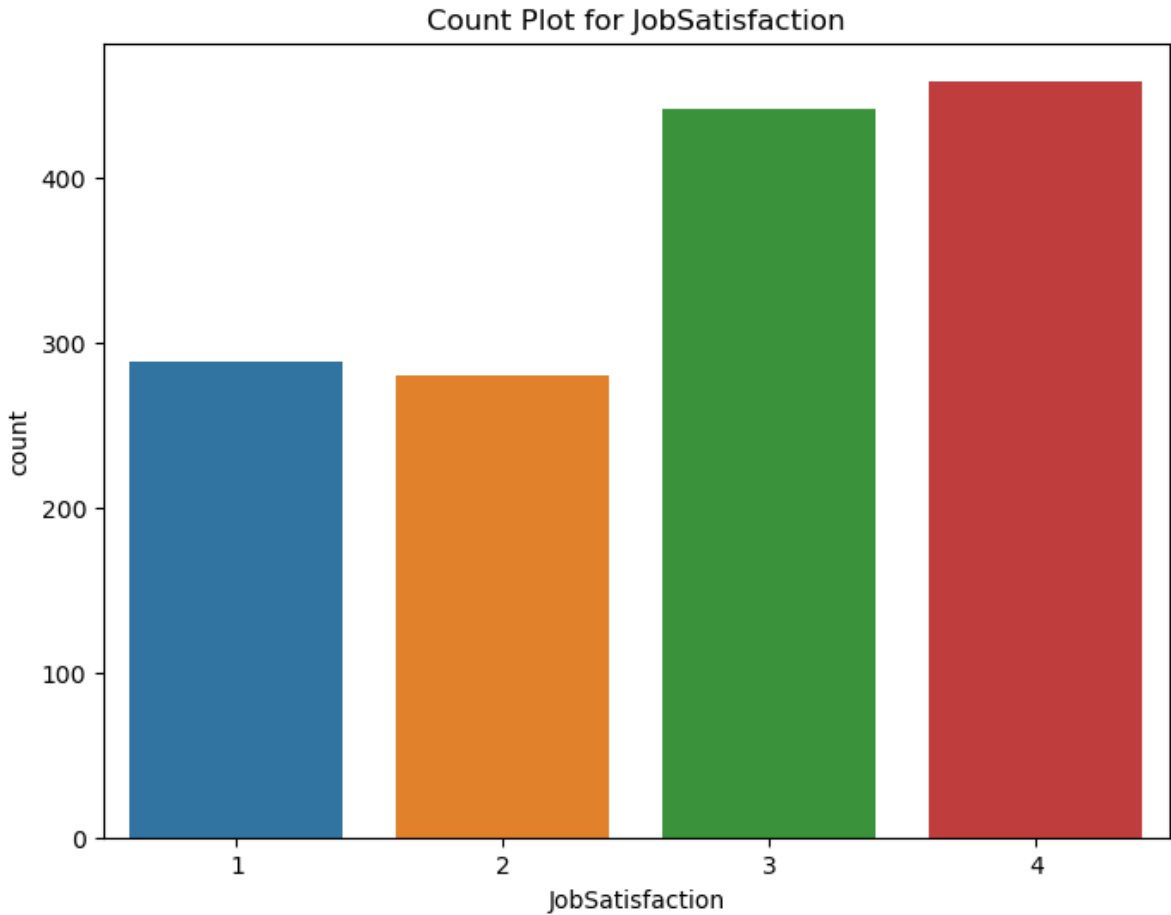
In [120...

```
#for the count plot:
selected_columns = ['EnvironmentSatisfaction', 'JobInvolvement', 'JobLevel', 'JobSatisfaction', 'NumCompaniesWorked', 'PercentSalaryHike', 'PerformanceRating', 'StandardHours', 'StockOptionLevel', 'TotalWorkingYears', 'TrainingPctCompleted', 'WorkLifeBalance', 'YearsAtCompany', 'YearsInCurrentRole', 'YearsWithCurrManager'] #selected_columns for the visualisation


for column in selected_columns:
    plt.figure(figsize=(8, 6))
    sns.countplot(x=column, data=employee_attrition)
    plt.title(f'Count Plot for {column}')
    plt.show()
```



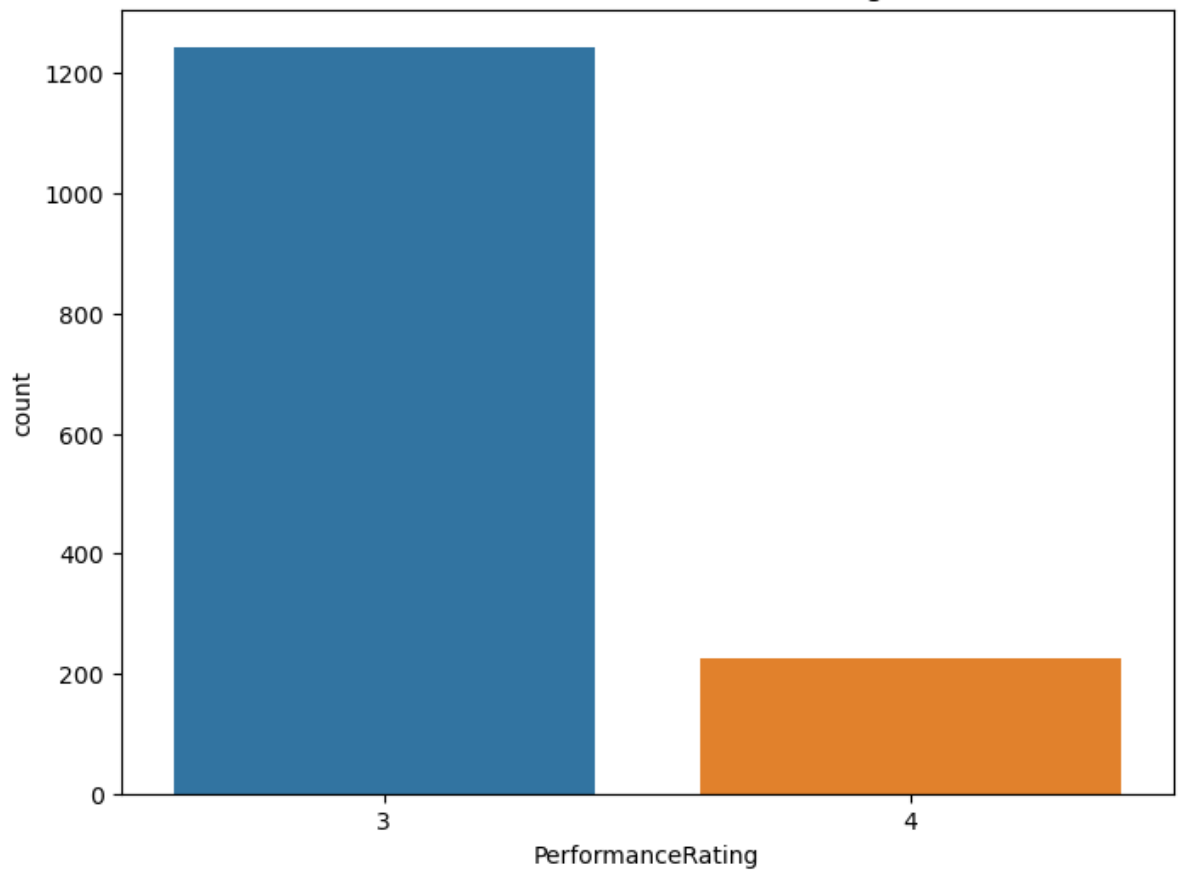


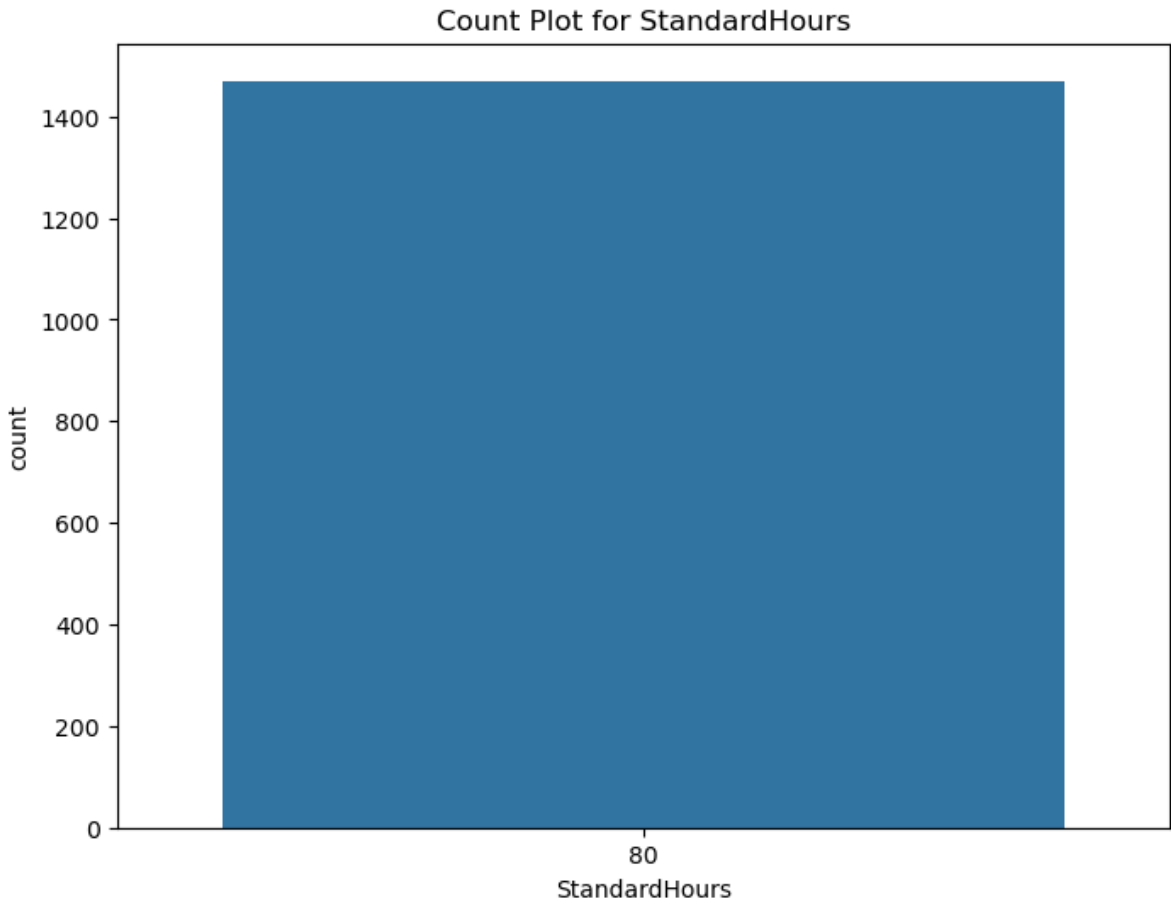
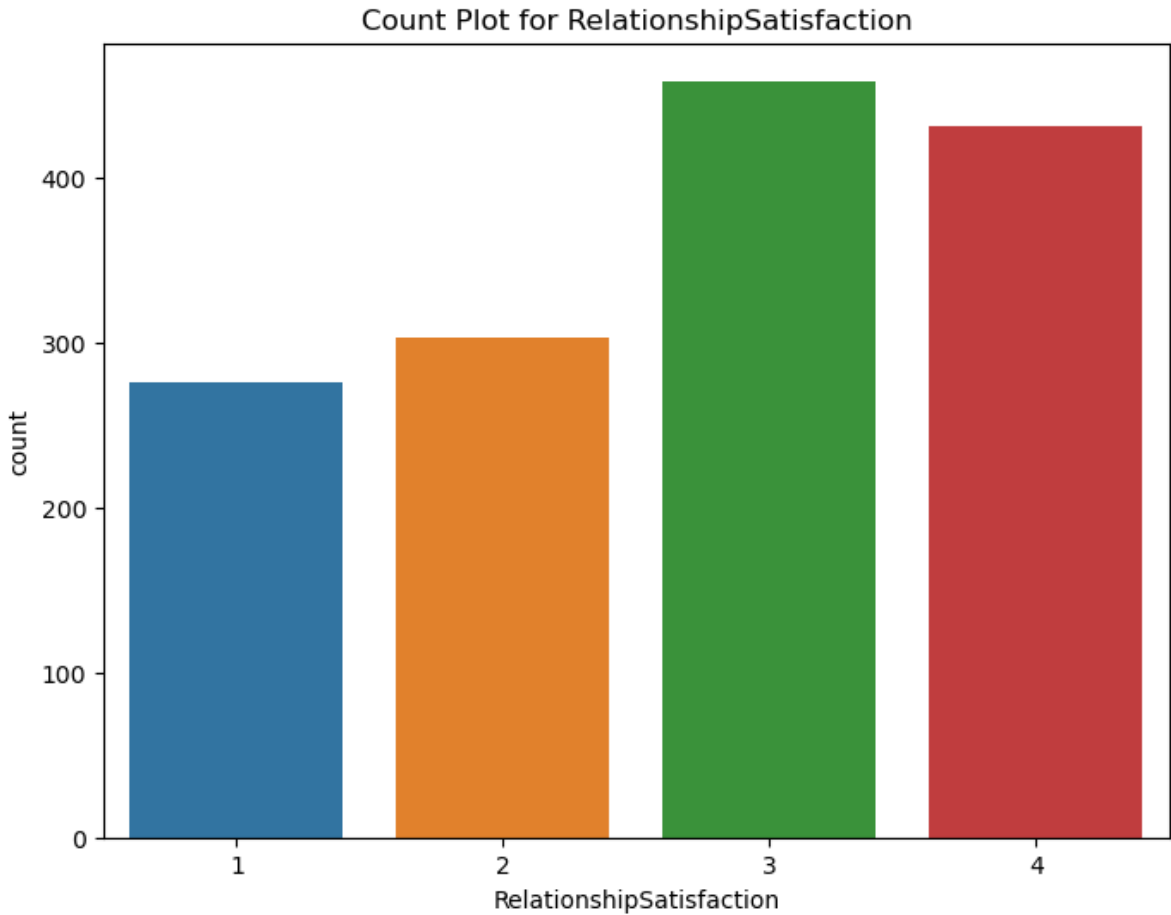


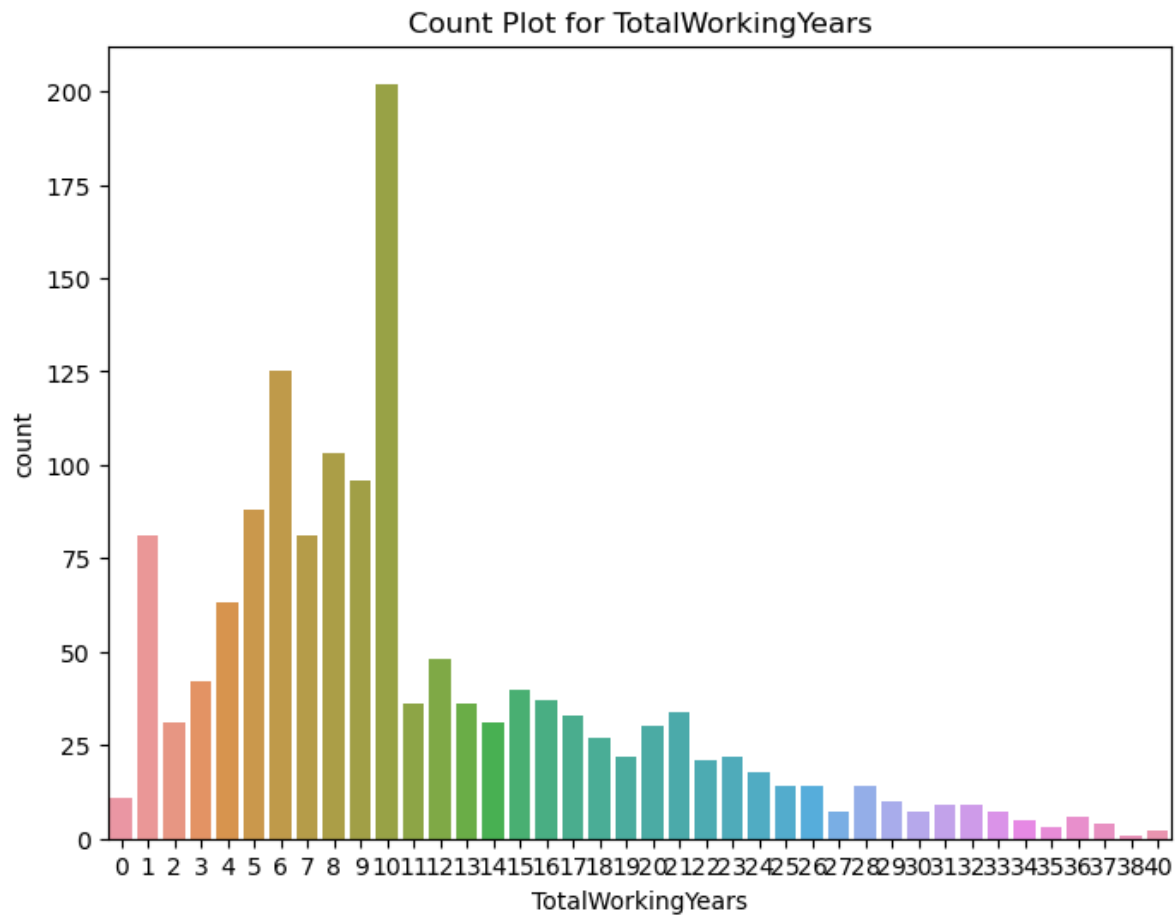
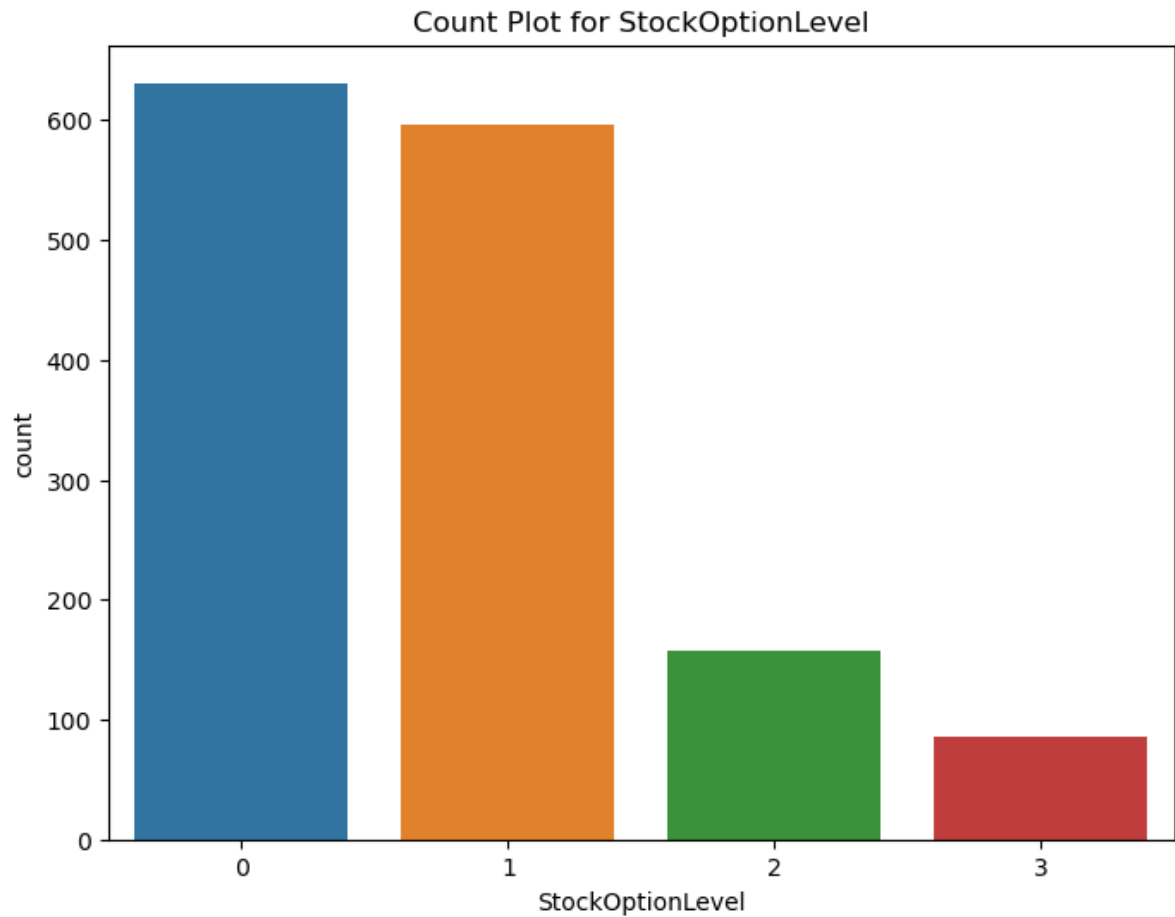
Count Plot for PercentSalaryHike

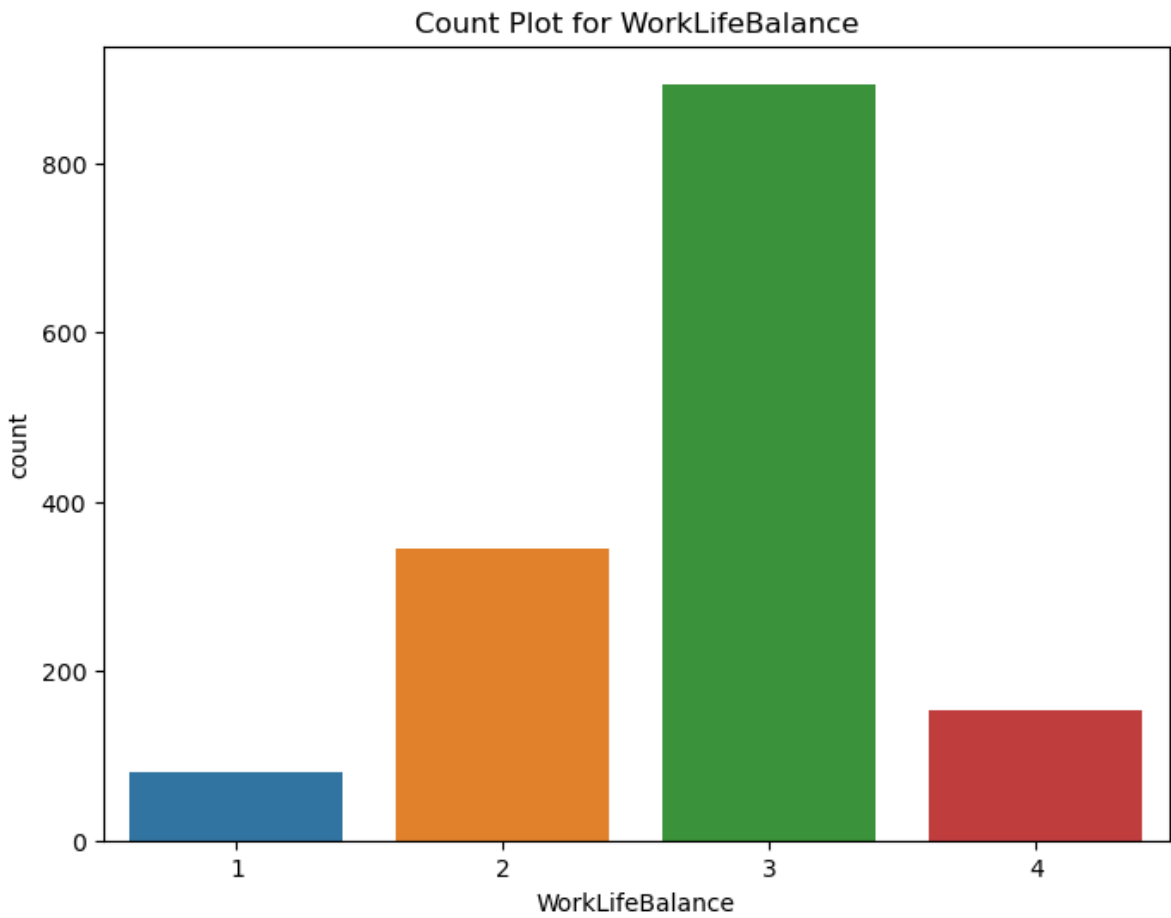
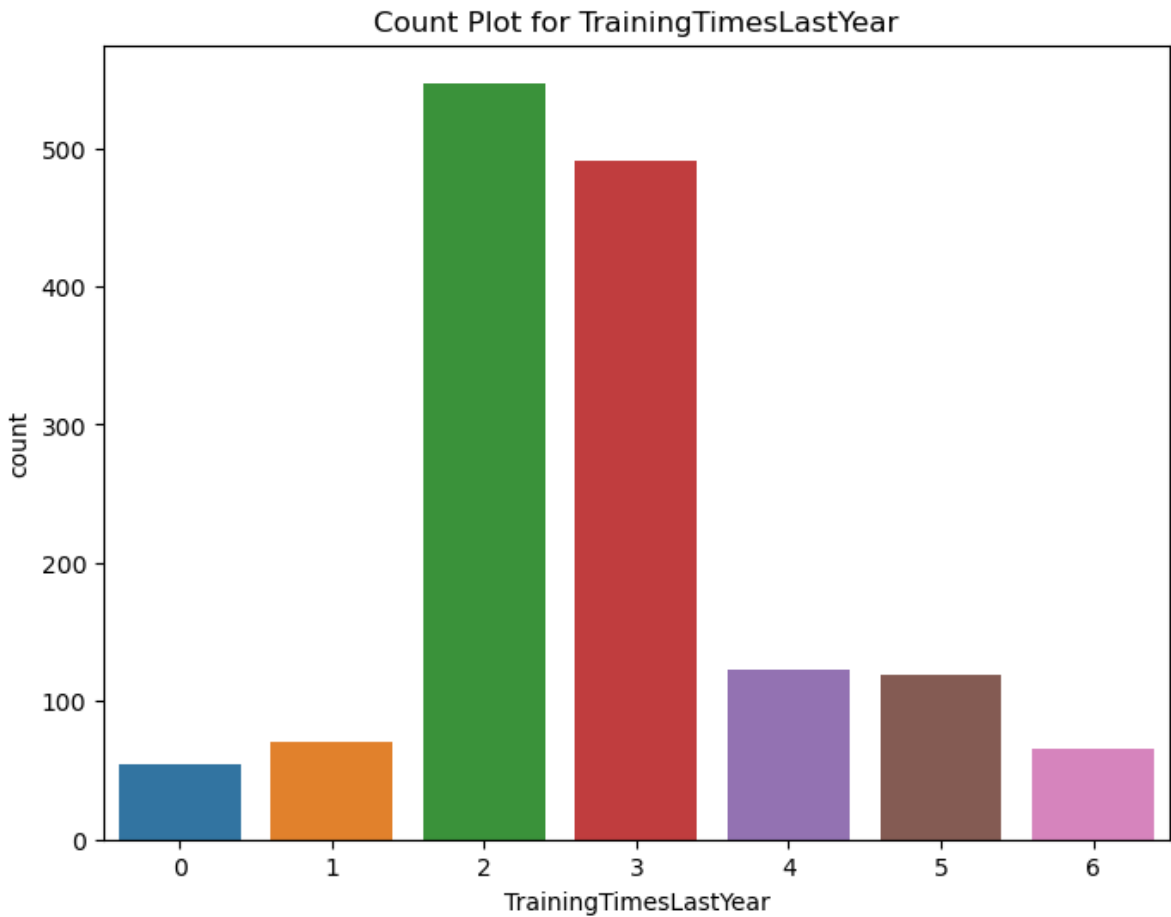


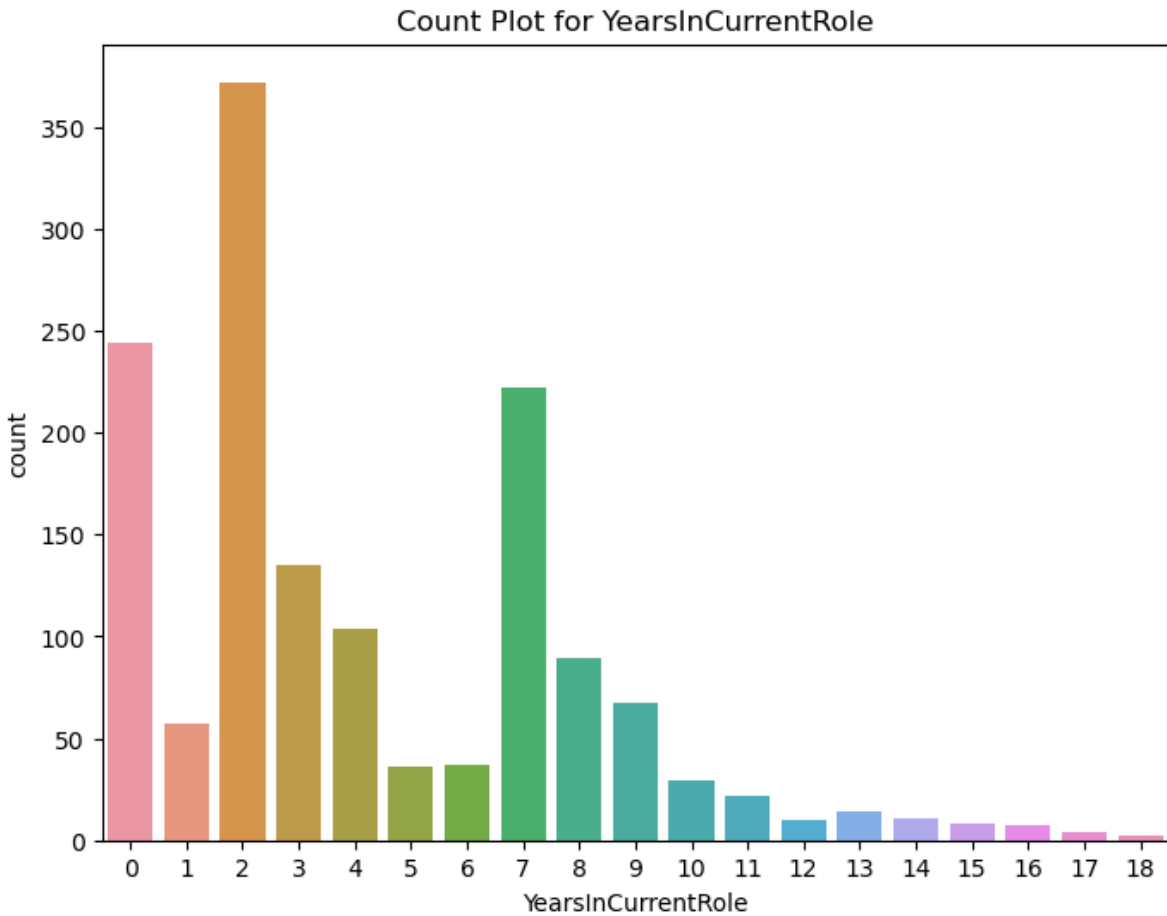

Count Plot for PerformanceRating


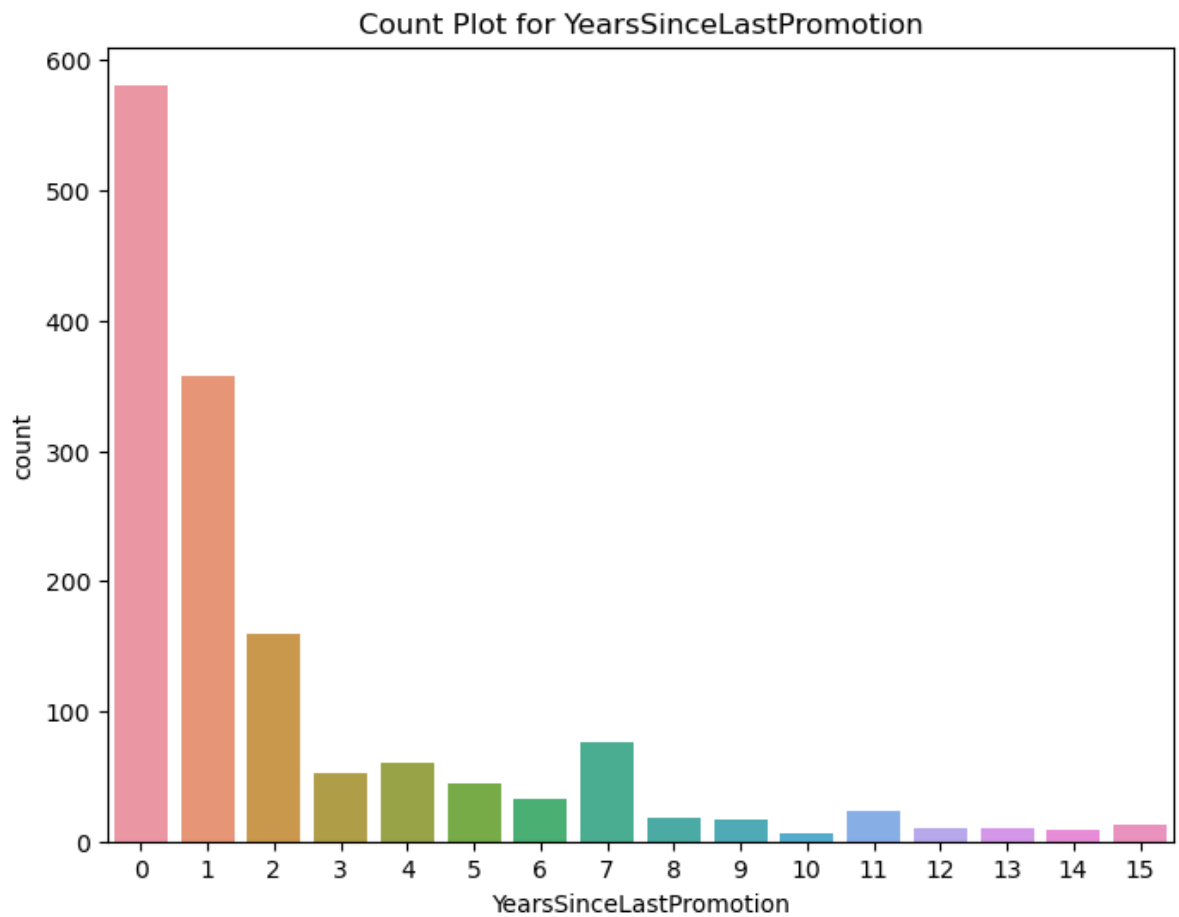










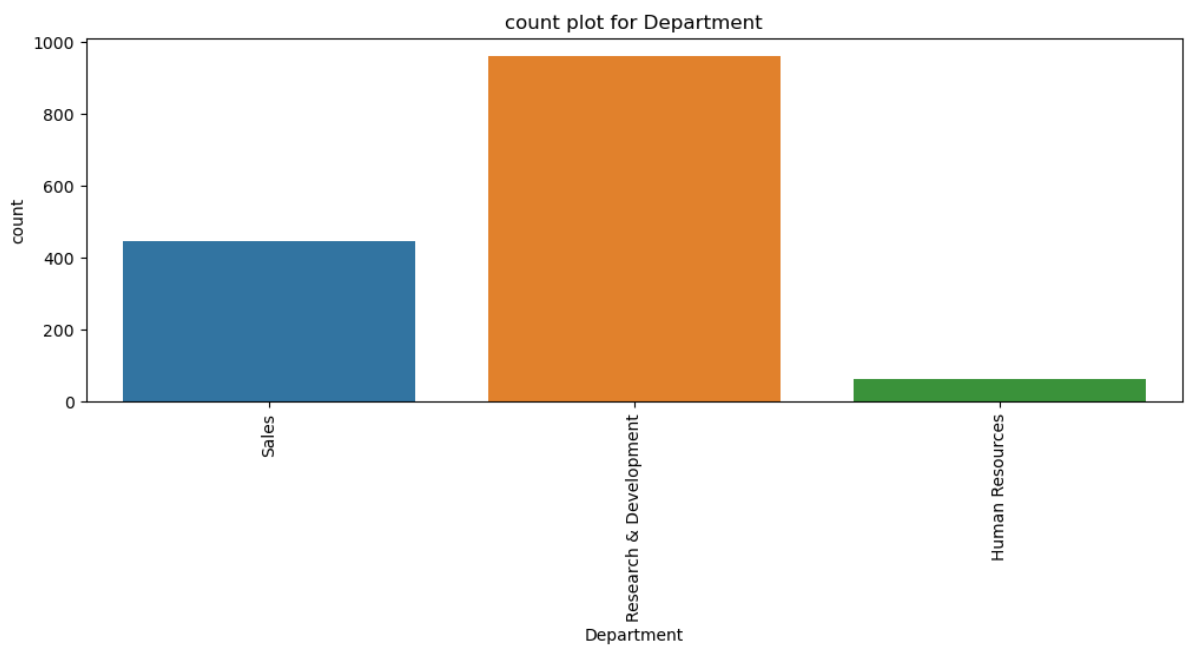
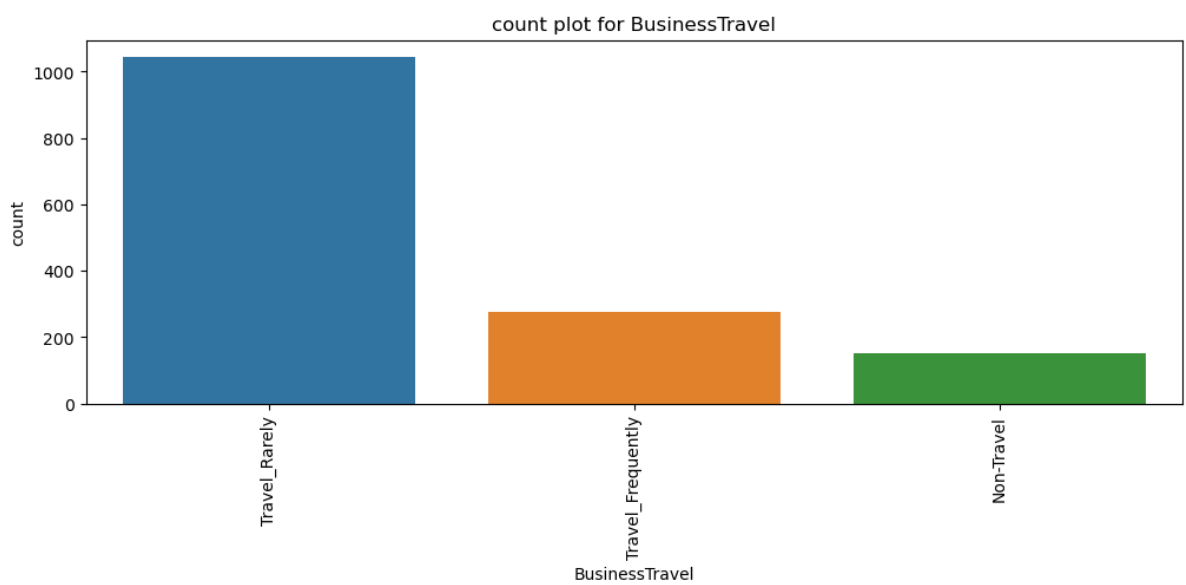
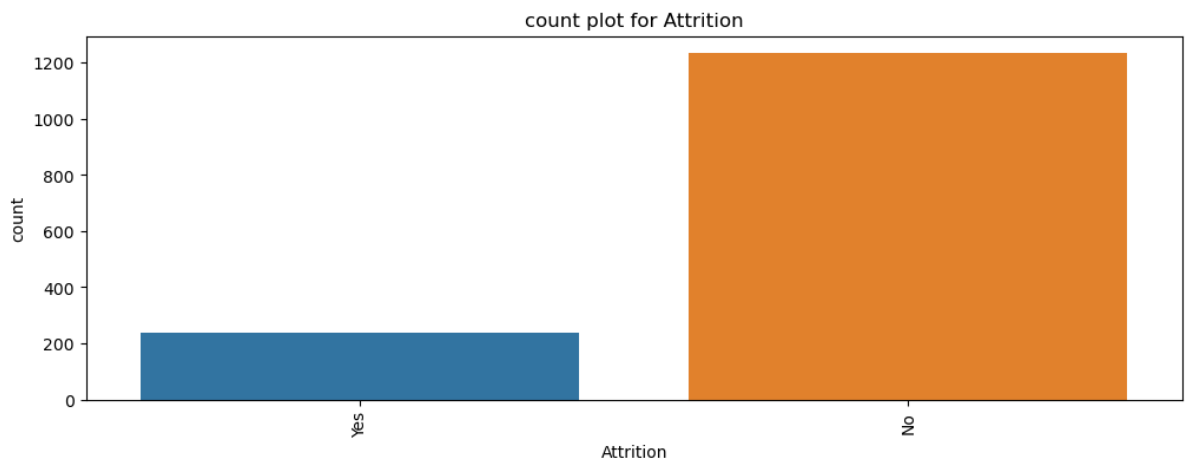


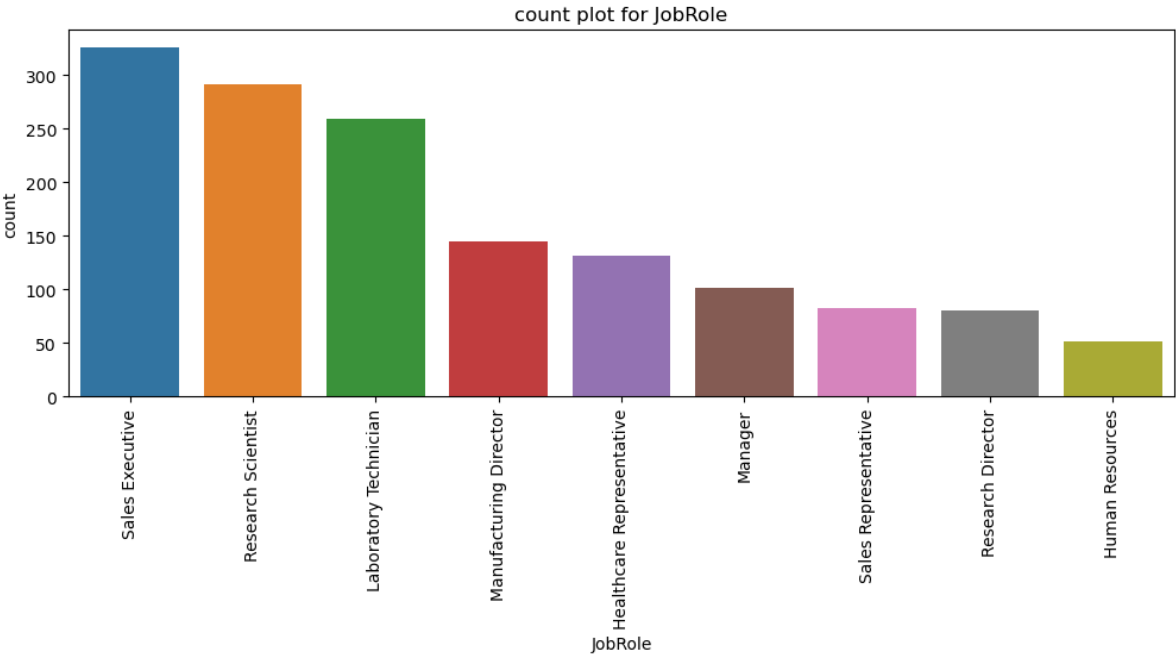

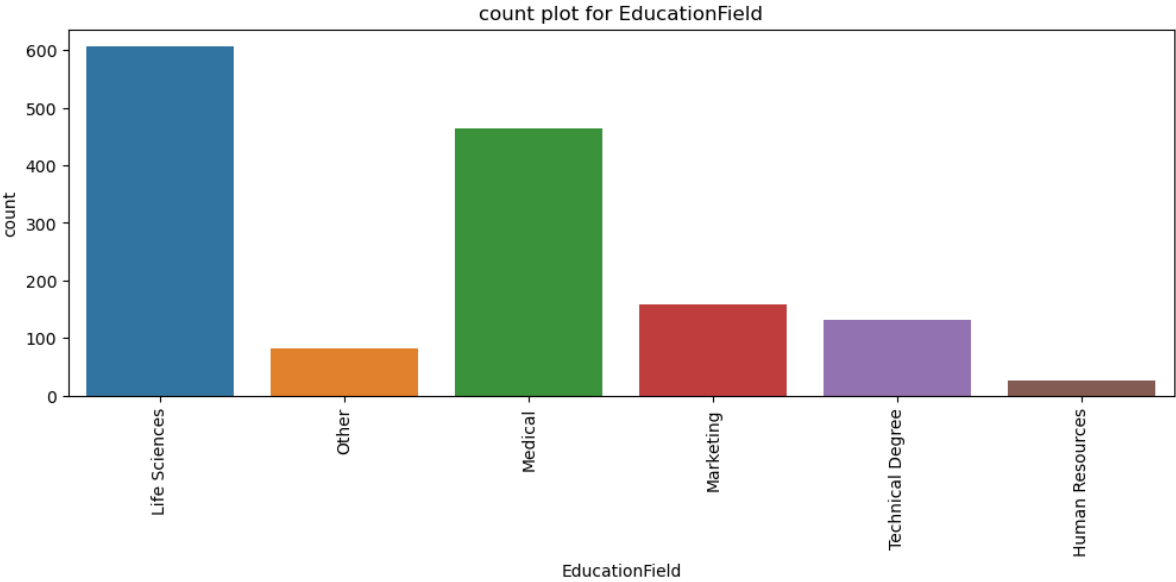
```
In [121... #from the info object, it is seen that the some data types are object. Therefore th

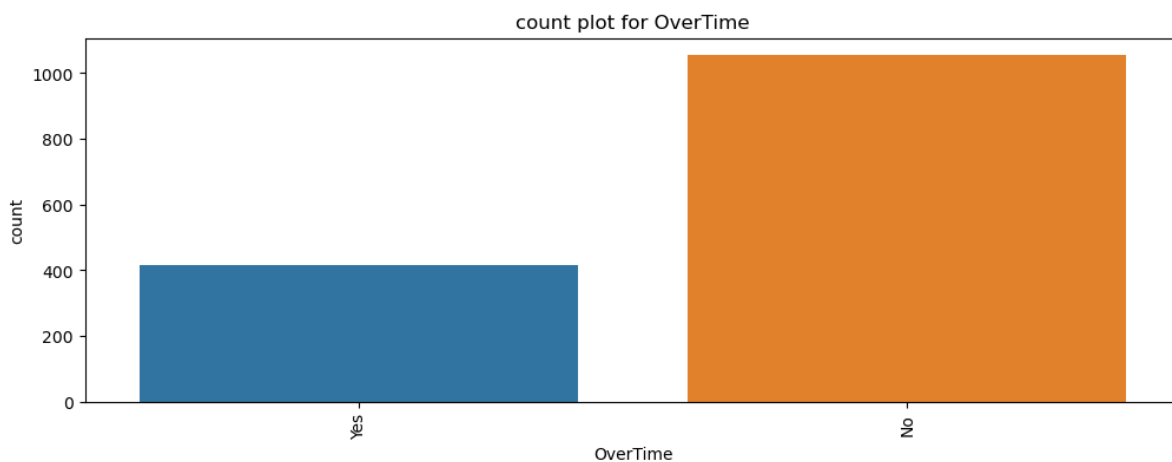
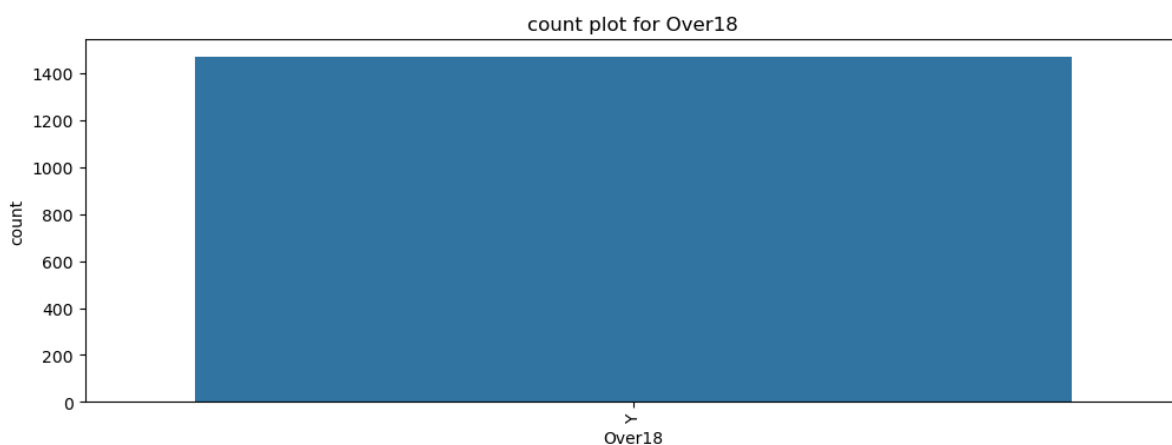

for column in employee_attrition.select_dtypes(include='object').columns:
    plt.figure(figsize=(12,4))
    plt.xticks(rotation=90)
```

```
sns.countplot(x=column, data=employee_attrition) #show the counts of observations  
plt.title(f'count plot for {column}')  
plt.show()
```

#here include='object' is used to represent string or categorical variables in pandas







```
In [122...] employee_attrition.columns
```

```
Out[122]: Index(['Age', 'Attrition', 'BusinessTravel', 'DailyRate', 'Department',
      'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount',
      'EmployeeNumber', 'EnvironmentSatisfaction', 'Gender', 'HourlyRate',
      'JobInvolvement', 'JobLevel', 'JobRole', 'JobSatisfaction',
      'MaritalStatus', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked',
      'Over18', 'OverTime', 'PercentSalaryHike', 'PerformanceRating',
      'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel',
      'TotalWorkingYears', 'TrainingTimesLastYear', 'WorkLifeBalance',
      'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotion',
      'YearsWithCurrManager'],
      dtype='object')
```

```
In [123...] object_columns = [col for col in employee_attrition.select_dtypes(include='object')]
```

```
In [124...] object_columns
```



```
Out[124]: ['Attrition',
           'BusinessTravel',
           'Department',
           'EducationField',
           'Gender',
           'JobRole',
           'MaritalStatus',
           'Over18',
           'OverTime']
```

```
In [125... selected_columns
```

```
Out[125]: ['EnvironmentSatisfaction',
           'JobInvolvement',
           'JobLevel',
           'JobSatisfaction',
           'NumCompaniesWorked',
           'PercentSalaryHike',
           'PerformanceRating',
           'RelationshipSatisfaction',
           'StandardHours',
           'StockOptionLevel',
           'TotalWorkingYears',
           'TrainingTimesLastYear',
           'WorkLifeBalance',
           'YearsAtCompany',
           'YearsInCurrentRole',
           'YearsSinceLastPromotion',
           'YearsWithCurrManager']
```


```
In [126... other_columns = []
for col in employee_attrition.columns:
    if col not in selected_columns and employee_attrition[col].dtype != 'object':
        other_columns.append(col) #we will add the rest of the columns to the other
```

```
In [127... other_columns #these columns were more suitable for boxplot representation
```


```
Out[127]: ['Age',
           'DailyRate',
           'DistanceFromHome',
           'Education',
           'EmployeeCount',
           'EmployeeNumber',
           'HourlyRate',
           'MonthlyIncome',
           'MonthlyRate']
```

```
In [128... for column in other_columns:
    sns.boxplot(x=column, data=employee_attrition)
    plt.title(f'Boxplot for {column}')
    plt.show()
```


Boxplot for Age




Boxplot for DailyRate




Boxplot for DistanceFromHome




Boxplot for Education




Boxplot for EmployeeCount




Boxplot for EmployeeNumber




Boxplot for HourlyRate



Boxplot for MonthlyIncome





Data Processing

- missing values

In [129...

```
#handling missing values  
missing_values = employee_attrition.isnull().sum()  
print(missing_values) # we can see that there are no missing values
```

```

Age                                0
Attrition                          0
BusinessTravel                     0
DailyRate                          0
Department                         0
DistanceFromHome                   0
Education                          0
EducationField                     0
EmployeeCount                      0
EmployeeNumber                     0
EnvironmentSatisfaction            0
Gender                             0
HourlyRate                         0
JobInvolvement                     0
JobLevel                           0
JobRole                            0
JobSatisfaction                    0
MaritalStatus                      0
MonthlyIncome                      0
MonthlyRate                        0
NumCompaniesWorked                 0
Over18                             0
OverTime                           0
PercentSalaryHike                  0
PerformanceRating                  0
RelationshipSatisfaction            0
StandardHours                      0
StockOptionLevel                   0
TotalWorkingYears                  0
TrainingTimesLastYear              0
WorkLifeBalance                    0
YearsAtCompany                     0
YearsInCurrentRole                 0
YearsSinceLastPromotion            0
YearsWithCurrManager               0
dtype: int64

```

- encoding categorical variables

```

In [130... # check the data types of each column
# previously the visualization is made possible by identifying data types from .inf
# but I will do it one more time

data_types = employee_attrition.dtypes
print(data_types)

```

```

Age                int64
Attrition          object
BusinessTravel     object
DailyRate         int64
Department        object
DistanceFromHome   int64
Education          int64
EducationField     object
EmployeeCount     int64
EmployeeNumber     int64
EnvironmentSatisfaction int64
Gender            object
HourlyRate        int64
JobInvolvement     int64
JobLevel          int64
JobRole           object
JobSatisfaction    int64
MaritalStatus     object
MonthlyIncome     int64
MonthlyRate       int64
NumCompaniesWorked int64
Over18            object
OverTime          object
PercentSalaryHike int64
PerformanceRating int64
RelationshipSatisfaction int64
StandardHours     int64
StockOptionLevel  int64
TotalWorkingYears int64
TrainingTimesLastYear int64
WorkLifeBalance   int64
YearsAtCompany    int64
YearsInCurrentRole int64
YearsSinceLastPromotion int64
YearsWithCurrManager int64
dtype: object

```

In [131...

```

#these types of datas should be interpreted as int as well
print(data_types[data_types == 'object'].index)

```

```

Index(['Attrition', 'BusinessTravel', 'Department', 'EducationField', 'Gender',
       'JobRole', 'MaritalStatus', 'Over18', 'OverTime'],
      dtype='object')

```

In [132...

```

print(data_types[data_types == 'int64'].index)

```

```

Index(['Age', 'DailyRate', 'DistanceFromHome', 'Education', 'EmployeeCount',
       'EmployeeNumber', 'EnvironmentSatisfaction', 'HourlyRate',
       'JobInvolvement', 'JobLevel', 'JobSatisfaction', 'MonthlyIncome',
       'MonthlyRate', 'NumCompaniesWorked', 'PercentSalaryHike',
       'PerformanceRating', 'RelationshipSatisfaction', 'StandardHours',
       'StockOptionLevel', 'TotalWorkingYears', 'TrainingTimesLastYear',
       'WorkLifeBalance', 'YearsAtCompany', 'YearsInCurrentRole',
       'YearsSinceLastPromotion', 'YearsWithCurrManager'],
      dtype='object')

```

In [133...

```

encoded_data = pd.get_dummies(employee_attrition, columns=object_columns, drop_first=True)
#Assume that all Yes are 1 and all 0 are No.
encoded_data.head()

```


Out[133]:

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	Environmen
0	41	1102	1	2	1	1	
1	49	279	8	1	1	2	
2	37	1373	2	2	1	4	
3	33	1392	3	4	1	5	
4	27	591	2	1	1	7	


5 rows × 48 columns

In [134...]

```
#Correlation plot
sns.heatmap(encoded_data.corr())
#from the figure the correlation seem to be mostly 0.
```

Out[134]:

<Axes: >



In [135...]

```
y = encoded_data['Attrition_Yes'].values #here is the y values in an array
X = encoded_data.drop('Attrition_Yes', axis=1).values
print(y)
print(X)
```

```
[1 0 1 ... 0 0 0]
[[ 41 1102 1 ... 0 1 1]
 [ 49 279 8 ... 1 0 0]
 [ 37 1373 2 ... 0 1 1]
 ...
 [ 27 155 4 ... 1 0 1]
 [ 49 1023 2 ... 1 0 0]
 [ 34 628 8 ... 1 0 0]]
```

- splitting the data as well as feature scaling with Standard Scaler

```
In [136... #splitting data into final test set and remaining data(64/16/20 split)

#First the test set will be seperated from the training set by a 20%

X_remain, X_test, y_remain, y_test = train_test_split(X, y, test_size = 0.2, random

print("X_test shape:" ,np.shape(X_test))
print("X_remain shape:" ,np.shape(X_remain))
print("y_test shape: ",np.shape(y_test))
print("y_remain shape:" ,np.shape(y_remain))

#After seperating train/test split, scaling is performed
#Fit on training data and transform training data

scaler = StandardScaler()
#fit_transform learns the transformation parameters from the training data and appl
X_train_scaled = scaler.fit_transform(X_remain)

#transform method applies the previously learned transformation to unseen data.
X_test_scaled = scaler.transform(X_test)

X_test shape: (294, 47)
X_remain shape: (1176, 47)
y_test shape: (294,)
y_remain shape: (1176,)
```

5-Fold Cross Validation

```
In [137... # Implementing 5-Fold Cross Validation
kf = KFold(n_splits = 5, shuffle = True, random_state = 42)
print(kf)

#for train_idx, cv_idx in kf.split(X_train_scaled, y_remain):
#    X_kf_train, X_kf_cv = X_train_scaled[train_idx], X_train_scaled[cv_idx]
#    y_kf_train, y_kf_cv = y_remain[train_idx], y_remain[cv_idx]

KFold(n_splits=5, random_state=42, shuffle=True)
```

K-Nearest Neighbors (KNN)

```
In [138... # first we can determine the k-value which represents the number of neighbors to co

#this part is from the codes we went over in the lectures.
def calculate_accuracy(y_true, y_pred):
    "calculating the accuracy of true and predicted labels"
    return np.mean(y_true == y_pred)

cv_scores = []

train_scores = []
val_scores = []
```

```

#Looking for different k values
for k in range(1,21):
    knn = KNeighborsClassifier(n_neighbors = k)

    fold_scores = []

    for train_idx, cv_idx in kf.split(X_train_scaled, y_remain):
        X_kf_train, X_kf_cv = X_train_scaled[train_idx], X_train_scaled[cv_idx]
        y_kf_train, y_kf_cv = y_remain[train_idx], y_remain[cv_idx]

        knn.fit(X_kf_train, y_kf_train)
        y_pred_cv = knn.predict(X_kf_cv)
        fold_scores.append(calculate_accuracy(y_kf_cv, y_pred_cv))

    #print("k:", k , "average CV accuracy for the corresponding k:" , np.mean(fold_scores))

    #taking the mean of the CV_results for the given k
    cv_scores.append(np.mean(fold_scores))

    #from the Lecture notes
    # Training the model on the entire training set. I did include the val and train
    knn.fit(X_train_scaled, y_remain)
    y_train_pred = knn.predict(X_train_scaled)
    train_scores.append(calculate_accuracy(y_remain, y_train_pred))

best_k = np.argmax(cv_scores) + 1
print("best_k: ", best_k)


# Plotting accuracies from the code in the Lecture notes
plt.figure(figsize=(10, 6))
plt.plot(range(1, 21), train_scores, label='Training Accuracy', marker='o')
plt.plot(range(1, 21), cv_scores, label='Cross-Validation Accuracy', marker='o')

# Highlighting the best k value from CV
plt.axvline(x=best_k, color='r', linestyle='--', label=f'Best k = {best_k}')

plt.title('k-NN Accuracies for Different k Values')
plt.xlabel('Value of k')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)
plt.show()

```

```
best_k: 11
```



```
In [139... # Now since we found the best_k we can test the accuracy now
best_knn = KNeighborsClassifier(n_neighbors = best_k) #best_k = 11

best_knn.fit(X_train_scaled, y_remain)

y_test_pred = best_knn.predict(X_test_scaled)

test_accuracy = calculate_accuracy(y_test, y_test_pred)

print(f"Test accuracy with best k: {k} is {test_accuracy:4f}")

Test accuracy with best k: 20 is 0.874150
```

Perceptron

```
In [140... fold_scores = []
for train_idx, cv_idx in kf.split(X_train_scaled, y_remain):
    X_kf_train, X_kf_cv = X_train_scaled[train_idx], X_train_scaled[cv_idx]
    y_kf_train, y_kf_cv = y_remain[train_idx], y_remain[cv_idx]

    perceptron = Perceptron()
    perceptron.fit(X_kf_train, y_kf_train)
    y_train_pred = perceptron.predict(X_kf_cv)
    fold_scores.append(calculate_accuracy(y_kf_cv, y_train_pred))

print(np.mean(fold_scores))
#From the 5-Fold Cross Validation the robustness of our model can be seen.

0.8180165885322754
```

```
In [141... # Create and train a perceptron model
perceptron = Perceptron()
perceptron.fit(X_train_scaled, y_remain)

# Make predictions on the test set
y_pred = perceptron.predict(X_test_scaled)

# Evaluate accuracy
```

```
accuracy = calculate_accuracy(y_test, y_pred)
print(f'Perceptron Accuracy: {accuracy}')
```

Perceptron Accuracy: 0.8639455782312925

Naive Bayes

In [145...]

```
#We may need to split the categorical and numerical features
#considering that the features are independent. We can add them later.
from sklearn.naive_bayes import GaussianNB

#GaussianNB() will be used
nb_Gaussian = GaussianNB()
nb_Gaussian.fit(X_train_scaled, y_remain)

# Make predictions
y_pred = nb_Gaussian.predict(X_test_scaled)

# Evaluate the model
accuracy = calculate_accuracy(y_test, y_pred)
print(f'Gaussian Navie Bayes Accuracy: {accuracy}')
```

Gaussian Navie Bayes Accuracy: 0.6904761904761905

Part II: Gradient Descent Implementation

1. The "Vehicle Dataset" should be downloaded from the Kaggle website:
<https://www.kaggle.com/datasets/nehalbirla/vehicle-dataset-from-cardekho>
2. Implement the gradient descent algorithm without using of any libraries except for Pandas and NumPy. [10 points]
3. How many iteration step needs to converge with learning rate [0.01, 0.1, 1]? Devise an intelligent strategy for choosing the learning rate to reduce the number of iterations required for convergence. Show how the learning rate that you propose impacts the convergence of the gradient descent algorithm. Show on the graph how the cost function changes with the number of iterations and how the gradient descent converges. [20 points]

In []:

```
# Your code here (you can add more blocks as you need).
# Please add comments where you think necessary.

#we since we seperated the
nb_categorical = Mu
```

Your Discussion Here (You can double click and edit this block)

In []:

```
#cite : https://towardsdatascience.com/what-is-feature-scaling-why-is-it-important-
```